

# Virtual Memory


# Background (Cont.)

- ❑ **Virtual memory** – separation of user logical memory from physical memory
  - ❑ Only part of the program needs to be in memory for execution
  - ❑ Logical address space can therefore be much larger than physical address space
  - ❑ Allows address spaces to be shared by several processes
  - ❑ Allows for more efficient process creation
  - ❑ More programs running concurrently
  - ❑ Less I/O needed to load or swap processes

# Background (Cont.)


- **Virtual address space** – logical view of how process is stored in memory
  - Usually start at address 0, contiguous addresses until end of space
  - Meanwhile, physical memory organized in page frames
  - MMU must map logical to physical
- Virtual memory can be implemented via:
  - Demand paging
  - Demand segmentation

# Virtual Memory That is Larger Than Physical Memory



# Demand Paging

- ❑ Could bring entire process into memory at load time
- ❑ Or bring a page into memory only when it is needed
  - ❑ Less I/O needed, no unnecessary I/O
  - ❑ Less memory needed
  - ❑ Faster response
  - ❑ More users
- ❑ Similar to paging system with swapping (diagram on right)
- ❑ Page is needed  $\Rightarrow$  reference to it
  - ❑ invalid reference  $\Rightarrow$  abort
  - ❑ not-in-memory  $\Rightarrow$  bring to memory
- ❑ **Lazy swapper** – never swaps a page into memory unless page will be needed
  - ❑ Swapper that deals with pages is a **pager**



# Basic Concepts

- With swapping, pager guesses which pages will be used before swapping out again
- Instead, pager brings in only those pages into memory
- How to determine that set of pages?
  - Need new MMU functionality to implement demand paging
- If pages needed are already **memory resident**
  - No difference from non demand-paging
- If page needed and not memory resident
  - Need to detect and load the page into memory from storage
    - ▶ Without changing program behavior
    - ▶ Without programmer needing to change code

# Valid-Invalid Bit

- With each page table entry a valid–invalid bit is associated (**v**  $\Rightarrow$  in-memory – **memory resident**, **i**  $\Rightarrow$  not-in-memory)
- Initially valid–invalid bit is set to **i** on all entries
- Example of a page table snapshot:

Frame #	valid-invalid bit
	<b>v</b>
	<b>v</b>
	<b>v</b>
	<b>i</b>
...	
	<b>i</b>
	<b>i</b>

page table

- During MMU address translation, if valid–invalid bit in page table entry is **i**  $\Rightarrow$  page fault

# Page Table When Some Pages Are Not in Main Memory

0	A
1	B
2	C
3	D
4	E
5	F
6	G
7	H


logical  
memory

	frame	valid-invalid bit
0	4	v
1		i
2	6	v
3		i
4		i
5	9	v
6		i
7		i

page table

0
1
2
3
4
5
6
7
8
9
10
11
12
13
14
15

physical memory





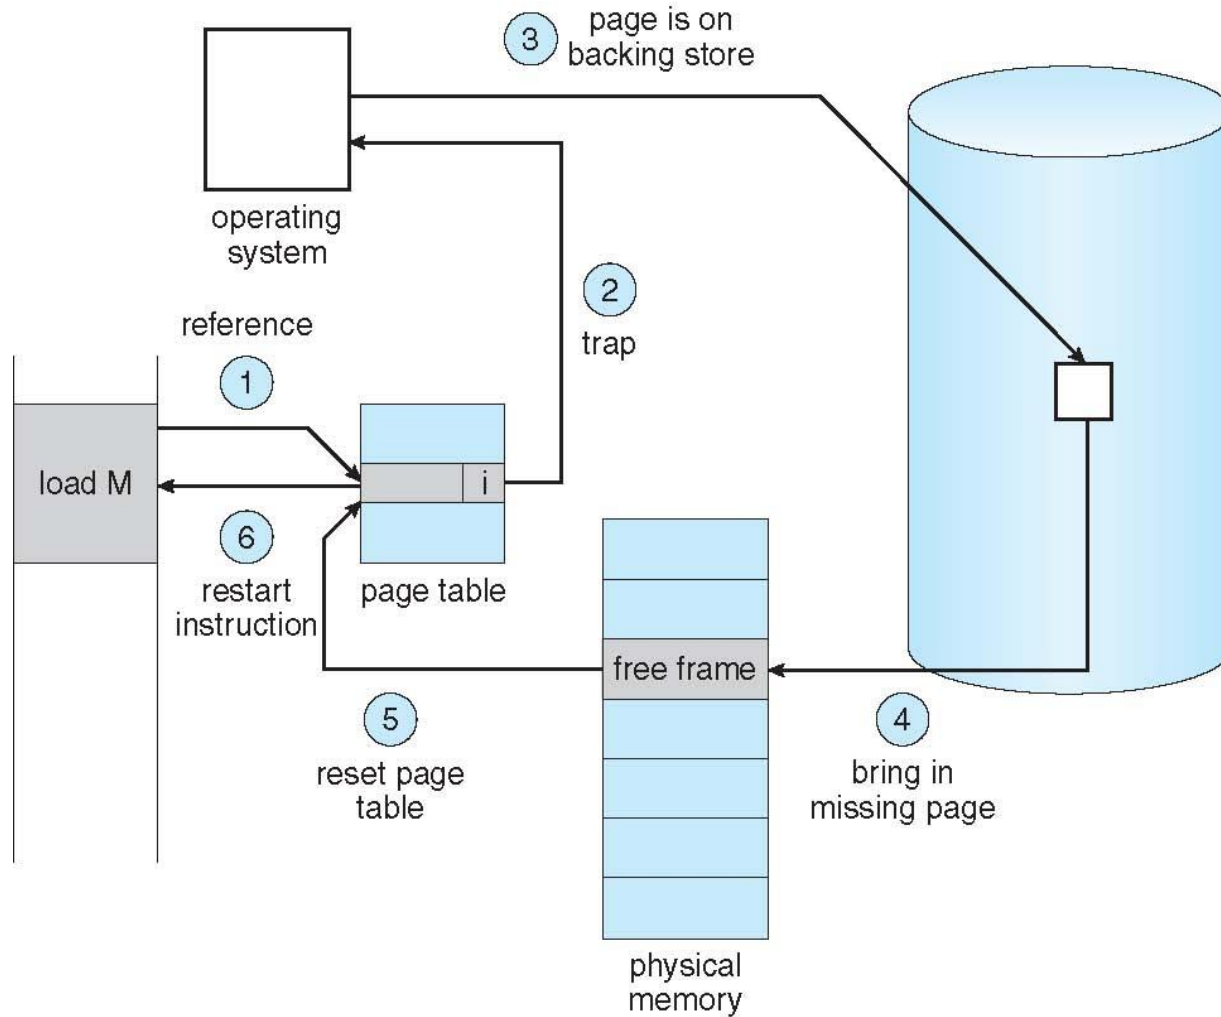
# Page Fault

- If there is a reference to a page, first reference to that page will trap to operating system:

## page fault

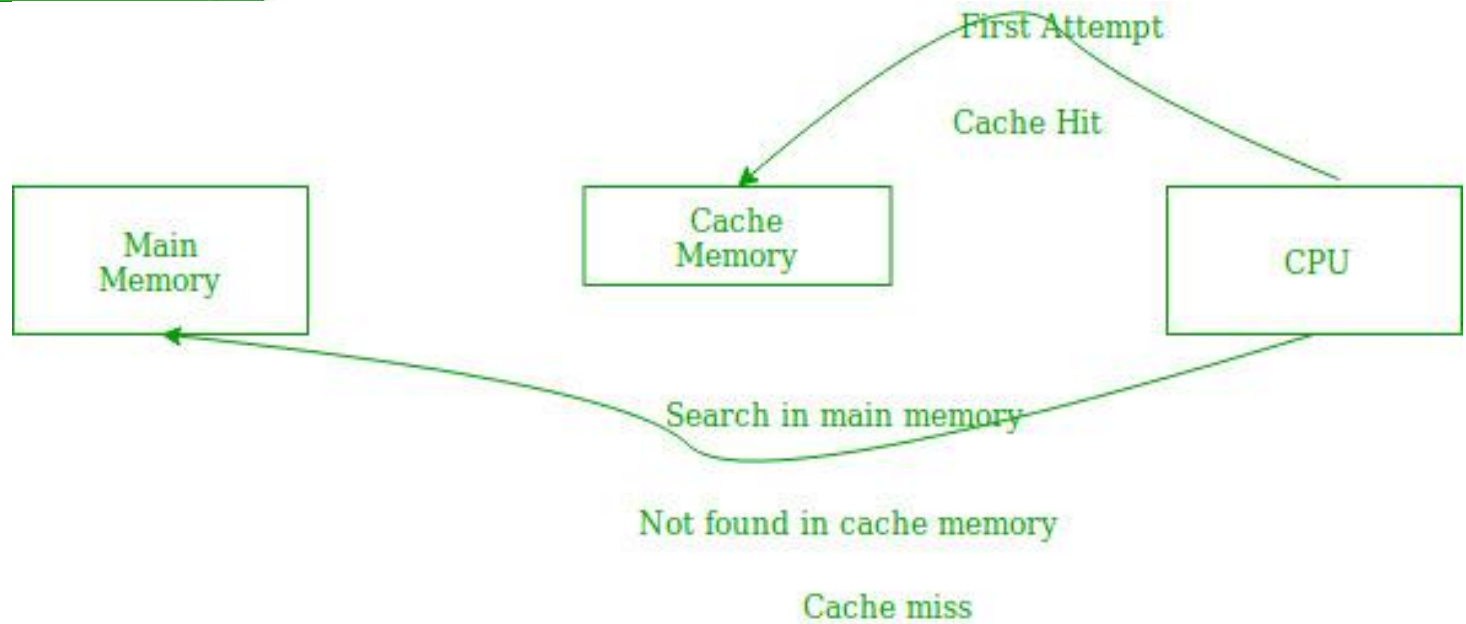
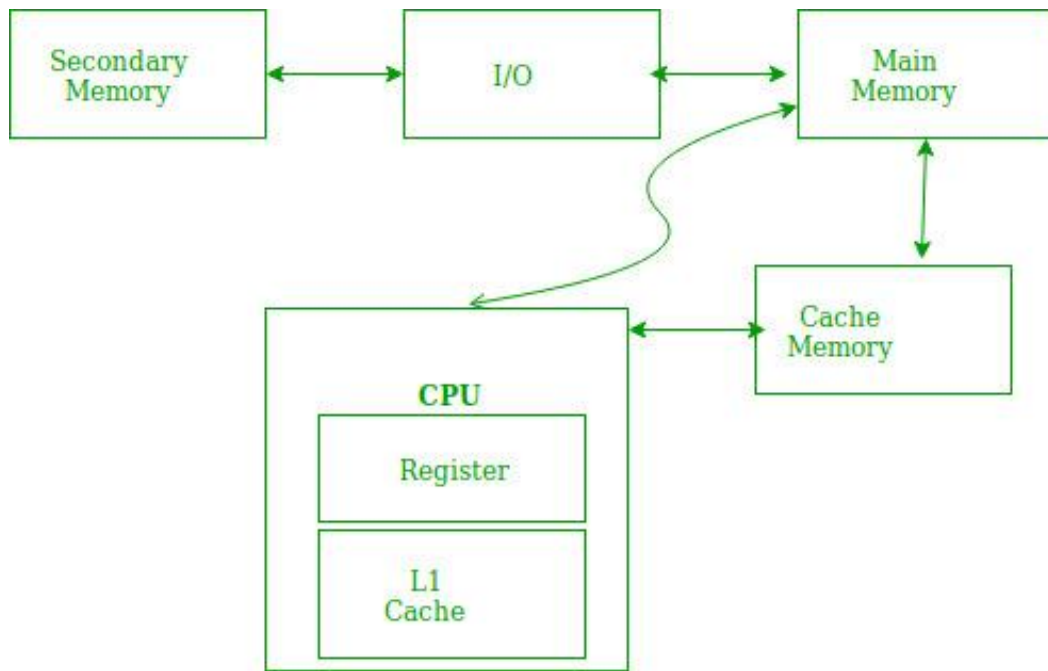
1. Operating system looks at another table to decide:
  - Invalid reference  $\Rightarrow$  abort
  - Just not in memory
2. Find free frame
3. Swap page into frame via scheduled disk operation
4. Reset tables to indicate page now in memory  
Set validation bit = **v**
5. Restart the instruction that caused the page fault

# Steps in Handling a Page Fault



# Aspects of Demand Paging


- ❑ Extreme case – start process with *no* pages in memory
  - ❑ OS sets instruction pointer to first instruction of process, non-memory-resident -> page fault
  - ❑ And for every other process pages on first access
  - ❑ **Pure demand paging**
- ❑ Actually, a given instruction could access multiple pages -> multiple page faults
  - ❑ Consider fetch and decode of instruction which adds 2 numbers from memory and stores result back to memory
  - ❑ Pain decreased because of **locality of reference**
- ❑ Hardware support needed for demand paging
  - ❑ Page table with valid / invalid bit
  - ❑ Secondary memory (swap device with **swap space**)
  - ❑ Instruction restart



# Page Replacement

- ❑ Prevent **over-allocation** of memory by modifying page-fault service routine to include page replacement
- ❑ Use **modify (dirty) bit** to reduce overhead of page transfers – only modified pages are written to disk
- ❑ Page replacement completes separation between logical memory and physical memory – large virtual memory can be provided on a smaller physical memory

# Need For Page Replacement



# Page and Frame Replacement Algorithms

- **Frame-allocation algorithm** determines
  - How many frames to give each process
  - Which frames to replace
- **Page-replacement algorithm**
  - Want lowest page-fault rate on both first access and re-access
- Evaluate algorithm by running it on a particular string of memory references (reference string) and computing the number of page faults on that string
  - String is just page numbers, not full addresses
  - Repeated access to the same page does not cause a page fault
  - Results depend on number of frames available
- In all our examples, the **reference string** of referenced page numbers is

**7,0,1,2,0,3,0,4,2,3,0,3,0,3,2,1,2,0,1,7,0,1**

# First-In-First-Out (FIFO) Algorithm

- Reference string: **7,0,1,2,0,3,0,4,2,3,0,3,0,3,2,1,2,0,1,7,0,1**
- 3 frames (3 pages can be in memory at a time per process)

reference string

7 0 1 2 0 3 0 4 2 3 0 3 0 3 2 1 2 0 1 7 0 1

7	7	7	2																		
	0	0	0																		
			1	1																	

page frames

15 page faults

- Can vary by reference string: consider 1,2,3,4,1,2,5,1,2,3,4,5
  - Adding more frames can cause more page faults!
    - **Belady's Anomaly**
- How to track ages of pages?
  - Just use a FIFO queue



# Optimal Algorithm

- Replace page that will not be used for longest period of time
  - 9 is optimal for the example
- How do you know this?
  - Can't read the future
- Used for measuring how well your algorithm performs

reference string

7 0 1 2 0 3 0 4 2 3 0 3 2 1 2 0 1 7 0 1

7	7	7	2		2		2		2		2						7		
	0	0	0		0		4		0		0						0		
		1	1		3		3		3		1						1		

page frames

# Least Recently Used (LRU) Algorithm

- Use past knowledge rather than future
- Replace page that has not been used in the most amount of time
- Associate time of last use with each page

reference string

7 0 1 2 0 3 0 4 2 3 0 3 2 1 2 0 1 7 0 1


7	7	7	2		2		4	4	4	0			1		1		1		
	0	0	0		0		0	0	3	3			3		0		0		
		1	1		3		3	2	2	2			2		2		7		

page frames

- 12 faults – better than FIFO but worse than OPT
- Generally good algorithm and frequently used
- But how to implement?

# Thrashing

- If a process does not have “enough” pages, the page-fault rate is very high
  - Page fault to get page
  - Replace existing frame
  - But quickly need replaced frame back
  - This leads to:
    - ▶ Low CPU utilization
    - ▶ Operating system thinking that it needs to increase the degree of multiprogramming
    - ▶ Another process added to the system
- **Thrashing** ≡ a process is busy swapping pages in and out



- What happens, if CPU want to access “Page 2” of each process
  - Increase in Page fault rate and most of the time will spent in bringing pages in and out of the RAM.
  - This process is referred to THRASHING

## Thrashing (Cont.)

